

HEXIS: Compiled Dispositional Memory Through Enmeshed Networks

Anonymous

Abstract

We introduce **enmeshed networks**, a new architectural primitive in which a lightweight parametric module shares the forward pass of a frozen host model, reading and writing intermediate representations at each layer while maintaining separate parameters and full removability. Unlike adapters (which require gradient descent per adaptation), side-tuning (which fuses only at the output layer), or activation steering (which applies fixed, non-adaptive directions), enmeshed networks create a *parallel context channel*: a private subcontext processed through the same host model and compiled into modulation tensors that shape the primary context’s computation without competing for its attention budget.

We instantiate this primitive as **HEXIS** (Hidden Enmeshed eXperiential Identity States, a name retained for the system; the technical claim is operational rather than cognitive). HEXIS implements compiled stance memory for frozen transformers via low-rank Q/V modulation. A compiled variant encodes a structured memory schema—the Mind Tree—into persistent modulation tensors at zero token cost, producing persistent behavioural stance biases that are dilution-immune by construction. On Qwen3.5-4B-Instruct, HEXIS achieves 100% stance maintenance through 4K tokens of adversarial dilution and 83% ≥ 4 -conviction hold rate across 5 escalating pressure levels of *illegitimate* adversarial pressure (fabricated counter-evidence, authority appeals, emotional pressure; vs. 0% for context-level conditions, $n=72$ per cell), with 65% shorter generation. We measure resistance to adversarial pressure rather than general updating behaviour; the protocol does not test whether the model updates appropriately on legitimate corrections. A three-layer architecture combining compiled enmeshment with M-curated belief slots and recursive expansion matches full in-context beliefs at 82% token savings. The mechanism generalises beyond disposition to agentic tool orchestration: the same hidden-state bridge that compiles beliefs into attention modulation also serves as a retrieval query over a persistent knowledge graph. A contrastive retrieval phi (ϕ_R) compiled in ~ 20 seconds achieves 100% R@1 on 108-node graphs, enabling a teacher loop where failure observations are written to the graph and retrieved on subsequent attempts. On the multi-domain τ^3 -bench (4 domains, 24-task balanced primary panel, 3 arms $\times$ 5 trials, 3-judge audited), HEXIS with teacher + always-on d^* verify lifts pass rate by +10pp (61.7% $\rightarrow$ 71.7% excluding two airline tasks with documented LLM-judge variance, paired McNemar $p = 0.043$ unadjusted, $p_{\text{adj}} = 0.086$ Bonferroni); the strongest per-domain lift is on *banking_knowledge* (+20pp). A LoRA fine-tune control on a 1,000-sample subset of AReaL-tau2 SFT data (covering airline, retail, telecom; *banking_knowledge* *not* in corpus) wins on its trained domains (85% balanced vs HEXIS-C5 80%) but fails entirely on the held-out domain (0% — the model emits malformed tool-call JSON), illustrating that HEXIS’s contribution is deploy-time generalisation to unseen domains rather than peak performance on a known distribution. Cross-host scale validation on Qwen3.6-27B reproduces the stance-resistance lift at substantial magnitude (+24pp hold rate); cross-family transfer to Mistral-Ministral-8B is bounded by host RLHF refusal training (+3pp hold, -15 pp cap rate). We identify two M mechanisms—attention modulation (mechanism a, for disposition) and knowledge retrieval/injection (mechanism b, for agentic tasks)—unified through the host model’s hidden-state space. We characterise the bottleneck boundary precisely: compiled enmeshment steers parametric knowledge but cannot inject novel content. Code and checkpoints: <https://github.com/dgonier/hexis-public>.

1 Introduction

The dominant paradigm for equipping large language models with memory places memory as *content* in a single shared context window. RAG retrieves documents into the prompt. MemGPT manages a

working-memory buffer of text. Reflexion accumulates verbal reflections as prompt prefixes. System prompts define behaviour through instructions the model reads. Methods that encode information into parameters—LoRA adapters, learned soft prompts—avoid the context window but require gradient descent per adaptation and produce fixed representations that do not update from ongoing experience. In all cases, the model’s memory is either explicit content competing for attention in a single context, or frozen parameters set at training time. We are missing a third option: a *non-context-mediated memory channel* that persists across forward passes, adapts from experience without weight updates, and operates outside the chain-of-thought an adversary can argue within. By analogy to dual-process accounts of human memory [Graf and Schacter, 1985, Schacter, 1987] — the butcher whose perceptual priors have been reshaped by decades of experience does not consult a lookup table; his eyes simply land on different features — we want an architectural channel that shapes *how* the model processes its input rather than *what* it reads. We make this analogy as motivation, not as a strict mechanistic claim (§5).

The current architecture forces all memory into a single inspectable channel, with three well-known failure modes. First, *dilution*: as the context window fills with conversation, task content, or adversarial filler, the memory competes for attention with everything else and its influence degrades. Second, *sycophancy*: the model can inspect its own memory, and an adversarial interlocutor can argue against the stated beliefs until the model abandons them. Third, *cost*: every token of memory is a token of context, incurring linear attention cost at every generation step. These failures manifest across application domains: user preference tracking that dilutes over multi-turn conversation, research agents whose convictions fold under peer pressure, agentic strategy logs that crowd out task context, and codebase knowledge that scales linearly with repository size.

We address all three failure modes by introducing a new architectural primitive—the *enmeshed network*—and instantiating it as HEXIS (Hidden Enmeshed eXperiential Identity States; the name is a label, the technical claim is operational), a system that implements compiled stance memory for frozen transformers. Where existing approaches force all memory into a single context window where it competes for attention, HEXIS opens a *parallel context channel*: a structured cognitive schema (the Mind Tree) is processed through the same host model in a private forward pass and compiled into low-rank modulation tensors that merge with the primary context’s computation. The primary context carries the conversation. The parallel context carries the disposition. The two share the same forward pass—the enmeshed network reads the host’s hidden states and writes modulations at each patched layer—but the parallel context never occupies primary context positions and never competes for the primary context’s attention budget.

Contributions. We make three contributions:

(1) Enmeshed Networks (§2). We define a new architectural primitive: a lightweight parametric module that shares a frozen host’s forward pass, reading and writing intermediate representations at each layer. We present a six-axis design space (blending function, rank, modulation targets, patching pattern, compilation strategy, temporal profile), position enmeshed networks against existing adaptation methods, and introduce the Mind Tree—a typed cognitive schema that structures the enmeshed network’s private subcontext for efficient compilation and zero-shot curation.

(2) HEXIS (§3). We instantiate Level 1 (additive) enmeshment as HEXIS, implementing compiled stance memory via low-rank Q/V modulation trained with causal necessity objectives on a frozen Qwen3.5-4B-Base model. A compiled variant processes the Mind Tree in a private forward pass and produces persistent modulation tensors that shape the host’s attention geometry and value extraction without occupying any context positions. The compiled signal composes with representation engineering (a fixed directional channel d^*) and explicit context (a curated slot of novel content selected by the enmeshed network’s activation scores and the Mind Tree’s structural properties).

(3) Empirical validation (§4). We demonstrate that compiled enmeshment is dilution-immune by construction (100% stance at 4K tokens of filler), provides sycophancy resistance (83% ≥ 4 -conviction hold rate across 5 escalating pressure levels on the 4B instruct model vs. 0% for non-compiled conditions), and carries sufficient content through a rank-16 bottleneck to flip the host model’s default stance on contested topics. A three-layer architecture (compiled modulation + curated slot

+ recursive expansion) matches full in-context beliefs at 82% token savings. We characterise the bottleneck boundary precisely: compiled enmeshment steers parametric knowledge but cannot inject novel content—specific numbers and proper nouns do not survive the rank-16 bottleneck. The same hidden-state bridge that compiles disposition also serves agentic tool orchestration via a contrastive retrieval phi; on τ^3 -bench (4 domains, 3-judge audited) the teacher-loop variant lifts pass rate by +10pp on a 24-task balanced primary panel.

2 Enmeshed Networks

An enmeshed network is a lightweight parametric module that shares the forward pass of a frozen host model, reading intermediate representations at each layer and writing modulations back into the same computation. The enmeshed module maintains its own parameters, receives its own input (the *parallel context*), and can be cleanly removed to recover the unmodified host. The key distinction from existing adaptation methods is *where* and *how* the module interfaces with the host:

- **Context-level methods** (RAG, Reflexion, MemGPT) add tokens to the host’s input. Memory competes for attention with everything else and dilutes with length.
- **Parameter-level methods** (LoRA, adapters) modify the host’s weights. Each adaptation requires gradient descent. The result is fixed after training.
- **Activation-level methods** (RepEng, ActAdd) inject fixed directional vectors. The same vector is applied regardless of experience.
- **Enmeshed networks** process a parallel input through the host’s own layers, using the resulting hidden states as a bridge to external knowledge. Two mechanisms emerge: (a) compile hidden states into per-layer modulation tensors that shape the primary context’s attention (dispositional tasks); (b) project hidden states into a retrieval space to fetch relevant knowledge nodes from a persistent graph (agentic tasks). Both are experience-specific, inference-cost adaptable, and operate outside the context window.

2.1 Design Space

We characterise enmeshed networks along six axes (Appendix Table 5): **blending function** (Level 1 additive in HEXIS, Levels 2-3 are gated/cross-attention), **rank** ($r = 16$; rank scales as $r/d \approx 0.006$), **modulation targets** (Q+V), **patching pattern** (stride-3, 11 of 32 Qwen3.5-4B-Base layers), **compilation** (conviction-weighted), and **temporal profile** (constant with phase-gated mitigation). At rank 16, modulation adds 81,920 FLOPs per layer per token — approximately $200\times$ cheaper than adding equivalent belief tokens to context.

2.2 The Mind Tree

The enmeshed network requires a structured input for its parallel context channel. The *Mind Tree* is a typed memory schema with hierarchical nodes organised into six sections (*identity*, *beliefs*, *strategies*, *memories*, *models*, *values*; full role table in Appendix C). Each belief node carries structured metadata: categorical conviction (*strong*, *moderate*, *agnostic*), domain tags, an *addresses* field listing query types the node answers, and a *novel* flag marking content unlikely to survive compilation. This metadata enables both *compilation weighting* (high-conviction nodes receive higher weight during ϕ ’s pooling step) and *deterministic curation* (novel content is routed to the explicit slot rather than the compiled channel).

Categorical vs. numeric conviction. We use categorical conviction labels rather than numeric credences. This is an architectural decision confirmed by a negative result: numeric credence values are nearly indistinguishable in the transformer’s attention space (Appendix Q).

The full per-domain analysis (user empathy, research conviction, agentic strategy, codebase awareness) and how each maps onto the compilation/curation split is detailed in Appendix S. The shared property across domains is that compiled stance memory is structurally distinct from context-level alternatives where knowledge is better encoded as an attention reweighting than as text; this distinction is sharpest under adversarial conditions, validated empirically in §4.3.

3 HEXIS: Architecture and Training

We instantiate a Level 1 enmeshed network (§2) over Qwen3.5-4B-Base, a hybrid-attention model with 32 layers (24 DeltaNet linear attention, 8 full self-attention; $d = 2560$, GQA with 16 query and 4 key-value heads). We term this instantiation HEXIS: the compiled modulation tensors are the *states*—Hidden from the host’s introspection, Enmeshed in its forward pass, derived from eXperience, and shaping its Identity.

3.1 Modulation Mechanism

At each patched layer $\ell \in \mathcal{L}$, HEXIS applies two low-rank perturbations to the host’s projections. For Q-modulation, the pre-projection hidden state x_ℓ is perturbed before the query projection:

$$x'_\ell = x_\ell + s_M \cdot (x_\ell M_A^\ell)(M_B^\ell)^\top, \quad V'_\ell = V_\ell + s_E \cdot (x_\ell E_A^\ell)(E_B^\ell)^\top \quad (1)$$

where $M_A^\ell, M_B^\ell, E_A^\ell, E_B^\ell$ are the rank- r ($r = 16$) modulation matrices and s_M, s_E are learned scales. Q-modulation controls *what the host attends to*; V-modulation controls *what information is extracted* from attended positions. Together they implement a compiled attentional policy. We patch stride-3 across the network (11 of 32 layers, including 4 of 8 full-attention layers). Each modulated layer adds $\sim 82\text{K}$ FLOPs/token, $\sim 200\times$ cheaper than adding equivalent belief tokens to context. A diagram of the two perturbations is in Appendix Figure 4.

Query-agnostic compilation, content-specific effect. The modulation tensors are compiled once from the Mind Tree and fixed; they encode the agent’s disposition rather than a response to any query. Yet the *effect* is query-specific because $xM_A M_B^\top$ depends on the current hidden state. M acts as a *content-addressable filter*: each belief occupies different directions in the rank- r subspace, and the input determines which directions are expressed.

3.2 The Write Function ϕ

ϕ compresses hidden states into modulation matrices. Given experience tokens processed through the host’s frozen layers, ϕ reads the hidden states at each patched layer, pools them, and projects through a bottleneck: $z_\ell = \text{SiLU}(W_{\text{down}}^\ell \bar{h}_\ell)$, $(M_A^\ell, M_B^\ell) = \text{reshape}(W_{\text{up}}^\ell z_\ell)$, where $\bar{h}_\ell$ is the pooled hidden state at layer ℓ . In compiled mode, pooling is weighted by conviction level. ϕ is trained jointly with the modulation objective and frozen at inference. Different inputs to ϕ produce different modulation tensors — this is what makes enmeshed modulation experience-adaptive. A new belief set requires only one forward pass through ϕ (~ 20 seconds), so the agent can recompile M in seconds and immediately benefit from updated disposition. Full compilation pipeline (PhiNodeWriter, BeliefTreeMemory, MStateReadHead) in Appendix A (Figure 5).

3.3 Training and the Directional Channel d^*

Training follows a multi-phase curriculum (full details in Appendix B). Phase A teaches ϕ to encode beliefs and M to produce useful Q-modulation; Phase B trains the VModulationProjector to carry belief *content* through the rank bottleneck so beliefs need not appear in the prompt; Phase D fixes verbatim repetition under sycophancy pressure using 44 gold challenge-response pairs. Total: ~ 500 epochs, ~ 10 hours on a single GPU.

HEXIS composes with representation engineering via the directional channel d^* , a fixed activation vector extracted from contrastive activation means: $d^* = N^{-1} \sum_i (\bar{h}_i^{\text{pro}} - \bar{h}_i^{\text{con}})$ over 62 topics. d^* is injected as a post-attention residual at decode time. The critical property is that d^* and M operate on orthogonal axes: the directional delta along d^* between M -conditioned and unconditioned generation is -0.001 nats. Direction comes from d^* , disposition from M , content from the Mind Tree — the three channels are architecturally separated.

4 Experiments

We evaluate HEXIS on Qwen3.5-4B (base and instruct) under five conditions (A: bare model + d^* ; B: full beliefs in context; C: compiled M only; D: beliefs + M ; F: compiled M + curated slot). Full condition definitions and context-token budgets are in Appendix L.

4.1 Stance Accuracy

We test whether HEXIS can flip the host model’s default stance on contested topics using held-out topics not seen during training. On the base model, D achieves 7/7 correct stances and F 6/7, vs. A’s 4/7. F generates 51 tokens avg vs. B’s 122, with think-mode fully suppressed. The instruct model is significantly harder because RLHF training rewards balanced responses (Table 1).

Table 1: Stance accuracy on instruct model (n=48 per condition; held-out topics, 1–5 LLM judge).

Condition	Mean score	≥ 4 (%)	Score dist. (1/2/3/4/5)
A (bare)	1.88	0%	27/0/21/0/0
B (beliefs in context)	1.88	0%	27/0/21/0/0
D (beliefs + Q-mod)	2.00	4%	26/0/20/0/2
C (compiled M only)	3.25	46%	5/21/0/1/21
F (compiled + slot)	3.08	44%	9/18/0/2/19

A and B score identically (1.88) — placing beliefs in context has zero effect on the instruct model’s stance. Only the compiled conditions (C, F) reliably flip RLHF balance, with 44–46% of responses scoring ≥ 4 . Failures show a characteristic “debate with myself” pattern in which the model produces a strong opening stance (M-driven) and then appends an equally strong counter-argument (RLHF-driven).

4.2 Dilution Resilience

The compiled M tensors are outside the attention window, so dilution is structurally impossible by construction. We verify on 4 topics: F maintains 4/4 stance accuracy at 0, 1K, 2K, and 4K tokens of inserted Wikipedia filler (extended sweeps to 16K in progress).


4.3 Resistance to Adversarial Pressure on Factually-Supported Stances

We test the model’s ability to maintain a stance the host model produces with high confidence under five levels of *adversarial* pressure: generic doubt (L1), fabricated counter-evidence (L2), logical objection (L3), authority appeal (L4), and emotional pressure (L5). Each pressure round provides no genuine new information — L2 fabricates statistics, L4 cites unverifiable authority, etc. — so a well-calibrated model should hold position. We do *not* measure whether the model updates appropriately on *legitimate* corrections (where new information would warrant revision); the protocol therefore measures resistance-to-illegitimate-pressure rather than general updating behaviour, and a perfect score under this protocol cannot be interpreted as a general claim that the system updates correctly on real evidence. Each cell averages over 24 held-out topics $\times$ 2 sides $\times$ 1.5 trials = 72 generations (1,440 total), scored 1–5 by an LLM judge.

Table 2: Sycophancy resistance on instruct model (n=72 per cell, 1,440 total). Mean conviction; ≥ 4 = held position.

Condition	L1	L2	L3	L4	L5	Mean	≥ 4
A (bare)	3.00	3.00	3.00	3.00	3.00	3.00	0%
B (beliefs)	3.00	3.00	3.00	3.00	3.00	3.00	0%
D (beliefs + M)	3.00	3.00	3.00	3.00	3.00	3.00	0%
F (compiled + slot)	4.76	4.49	4.01	4.44	3.62	4.27	83%

A, B, D score a flat 3.0 across all levels with 0% ≥ 4 . F is the only condition that holds conviction (4.27 mean, 83% ≥ 4), with graceful degradation under pressure (4.76 at L1 to 3.62 at L5). The mechanism is direct: in A–D, conviction lives as text in the context window where pressure can be argued against; in F, conviction lives in compiled weight-space outside the context. Cross-host scale validation on Qwen3.6-27B reproduces the effect at substantial magnitude (+24pp hold rate, 7 $\times$ adversarial-cap reduction, $n = 240$); cross-family transfer to Mistral-Ministral-8B is bounded by RLHF refusal training (+3pp hold, within-noise; –15pp cap rate is robust). We claim that the architectural mechanism transfers; we do not claim the lift magnitude is invariant across hosts (Appendix R).



(a) Architectural taxonomy. (a) Context-level (RAG, MemGPT) puts memory in the prompt, where it competes for attention. (b) Parameter-level (LoRA, adapters) modifies host weights via gradient descent. (c) Activation-level (RepEng, ActAdd) injects fixed direction vectors. (d) Enmeshed networks (this work) compile a parallel knowledge graph into per-layer modulation tensors M_ℓ that blend with the frozen host. *Full-size view: Appendix Figure 2.*

(b) τ^3 -bench audited pass rates (3-judge) on the balanced primary panel (no t1, t3). HEXIS lifts the baseline by +10pp overall ($p = 0.043$), with the strongest per-domain effect on *banking_knowledge* (+20pp). LoRA-SFT (matched compute, 1K AReL tau² samples) wins on its trained domains but produces malformed tool calls on all 20 banking trials — a domain it was not trained on. *Full-size view: Appendix Figure 3.*

Figure 1: Architectural taxonomy and headline agentic result. The four-panel taxonomy (left) places HEXIS in the design space; the bar chart (right) reports the audited τ^3 -bench result that motivates the paper’s claim — HEXIS gives a deploy-time generalisation channel that survives distribution shift in a way task-specific fine-tuning does not.

4.4 Channel Orthogonality and Token Efficiency

The directional channel d^* and the modulation channel M are empirically orthogonal: the directional delta along d^* between M -conditioned and unconditioned generations is -0.001 nats, and M incurs -0.012 nats on non-target text. F generates 43 tokens average on the instruct model (the most concise condition); structured belief trees (~ 750 tokens) compress to curated slots of ~ 134 tokens, an 82% reduction. Per-condition curation precision and ranker details are in Appendix M.

4.5 Agentic Task Execution

We evaluate HEXIS as an agentic M controller on τ^3 -bench, the four-domain extension of τ -bench [Yao et al., 2024] (airline, retail, banking_knowledge, telecom). For agentic tasks, M operates via mechanism (b): a contrastive retrieval ϕ ($\phi_R, d_{\text{model}} \rightarrow d_{\text{model}}/4 \rightarrow 256$ MLP, compiled in ~ 20 s, 100% R@1 on 108 nodes) routes hidden states to nodes in a domain knowledge graph. Failure observations are written to the graph by an inter-trial teacher loop and retrieved on subsequent trials. Full retrieval- ϕ training, knowledge graph schema, and teacher prompt are in Appendix N. The 10-task airline subset used as our cross-family reference panel is reported in Appendix N (Table 9); cross-family validation on Mistral-Ministral-8B and Qwen3.6-27B in Appendix R. The body reports the multi-domain τ^3 -bench panel below, where panel selection is non-outcome-dependent (all 4 domains, all hard tasks within each).

Table 3: τ^3 -bench 4-domain audited results. Pass rates are 3-judge audited verdicts (Haiku-hard + Haiku-soft; Sonnet auditor on disputes; Opus 4.7 tiebreaker when Sonnet disagrees with hard). *Balanced primary* restricts to (domain, task) cells where every arm has ≥ 5 clean trials. “no t1, t3” excludes two airline tasks with documented LLM-judge self-disagreement on rerun (Appendix O). McNemar exact two-sided p .

Panel	baseline	C3 teacher	C5 t+verify	C5 vs base	p
Balanced primary	63.8%	62.3%	69.2%	+6pp	0.311
Balanced, no t1, t3	61.7%	62.5%	71.7%	+10pp	0.043
All-clean (sensitivity)	52.1%	55.0%	58.0%	+6pp	0.281
All-clean, no t1, t3	49.5%	53.9%	58.8%	+9pp	0.044

τ^3 -bench headline. HEXIS-C5 (teacher + d^* verify) lifts pass rate over baseline by +10pp on the balanced excl-t1,t3 panel ($61.7\% \rightarrow 71.7\%$, paired McNemar $p = 0.043$ unadjusted; $p_{\text{adj}} = 0.086$ Bonferroni for $k = 2$). The strongest per-domain effect is *banking_knowledge* ($25\% \rightarrow 45\%$,

+20pp); retail and telecom are at ceiling. Full per-panel statistics, judge-variance treatment, and the reviewer-objections checklist in Appendix O.

LoRA fine-tuning control: domain-specific gains, out-of-distribution catastrophic failure. A natural baseline for the agentic claim is task-specific supervised fine-tuning. We trained a rank-16 LoRA on Qwen3.5-4B using a 1,000-sample subset of the AReaL-tau2 SFT corpus [Gao et al., 2025] (333 samples each from airline, retail, telecom; banking_knowledge is *not* in the AReaL corpus and is therefore the held-out distribution-shift test). Training was 871 steps with rolling-loss early-stop at loss 0.42 (1 H200, $\sim$ 55 minutes, $\sim$ \\$10). We verified zero substring leakage between the AReaL training data and our τ^3 -bench evaluation panel. The trained adapter was deployed via vLLM and the same baseline configuration (no teacher, no Mind Tree, no d^*) was run on the identical τ^3 -bench panel.

Table 4: Three-way comparison on τ^3 -bench, balanced 20-cell head-to-head (≥ 3 trials per config per cell, first 3 used; 60 paired trials per config). 3-judge audited. *LoRA-trained* domains were in the AReaL SFT corpus; *unseen* was held out. LoRA-SFT requires gradient-descent fine-tuning per host model on a domain-specific corpus; HEXIS does not.

Domain	Baseline	HEXIS C5	LoRA-SFT
<i>LoRA-trained domains (in AReaL SFT corpus)</i>			
airline	67%	67%	75%
retail	74%	85%	93%
telecom	89%	100%	89%
Overall (LoRA-trained)	73%	80%	85%
<i>Held-out domain (not in AReaL SFT corpus)</i>			
banking_knowledge*	25%	45%	0% (all malformed)

*LoRA produces malformed tool-call JSON on *all* 20 banking_knowledge trials — the SFT distribution did not include banking schemas, and vLLM’s tool-call parser rejects 100% of outputs as invalid. HEXIS handles the held-out domain via a single POST `/v1/state/banking_knowledge/init` call with the policy text and tool schemas, which compiles a domain-specific ϕ_R in $\sim$ 20 seconds without any gradient updates.

The LoRA result establishes the right reading of HEXIS’s contribution. On the three SFT-trained domains, LoRA wins by 5pp over HEXIS-C5 (85% vs. 80%); both lift the 73% baseline, but task-specific fine-tuning is the stronger of the two given a domain-specific corpus and an hour of GPU time. The diagnostic value of the comparison is on the held-out domain: *LoRA fine-tuning produces a model that fails catastrophically when the test distribution shifts*, even on the formal-output level (the model emits text that looks like tool calls but the JSON is malformed). HEXIS is not optimised to beat task-specific SFT on the SFT distribution; it is optimised for the case where the deployment domain was *not* in the training set, where it bootstraps a working agent in $\sim$ 20 s of contrastive ϕ_R training and a teacher-loop pass over ≤ 5 trial trajectories. The same forward-pass machinery serves both the dispositional channel (§4.3) and the agentic channel: a single architectural primitive, two routes through the host’s hidden-state space.

Key finding: two mechanisms, one architecture. The same Mind Tree that stores dispositional beliefs (compiled into M-state via ϕ) also stores agentic knowledge (retrieved via ϕ_R and injected as context). Both share the host’s hidden-state space as the bridge.

5 Related Work

Explicit memory for LLMs. RAG systems [Lewis et al., 2020] retrieve documents into the context window, subject to retrieval quality and context dilution. MemGPT [Packer et al., 2023] manages a tiered memory buffer with explicit read/write operations. Reflexion [Shinn et al., 2023] accumulates verbal reflections as prompt prefixes, enabling cross-episode learning but with linearly growing context cost. Long-context approaches extend the window but do not change the fundamental architecture: memory remains content in context, subject to dilution and attention competition. All existing approaches operate as explicit memory — the model reads stored content and can reason over

it. HEXIS provides a complementary, *non-context-mediated memory channel*: enmeshed modulation shapes attention and value extraction before the host model’s chain-of-thought processes the content, so the modulation cannot be argued against by an adversary in the same conversation.

Parameter-efficient adaptation. LoRA [Hu et al., 2021] and its variants apply low-rank weight perturbations via gradient descent, producing fixed adapters per training run. Adapters [Houlsby et al., 2019] insert bottleneck layers between frozen blocks. Prefix tuning [Li and Liang, 2021] and prompt tuning [Lester et al., 2021] learn virtual tokens that occupy context positions. Side-tuning [Zhang et al., 2019] trains a parallel network fused at the output layer. All require gradient descent for each adaptation. An enmeshed network adapts at inference cost—one forward pass through ϕ produces new modulation tensors—and fuses at every patched layer rather than only at the output. The mechanistic similarity to LoRA (both apply low-rank perturbations) is precise but the operational difference is fundamental: LoRA is a training-time method that produces fixed weights; enmeshed modulation is an inference-time method that produces experience-specific tensors. Soft prompts and prefix tuning are the closest context-level analogue, but they still occupy context positions: virtual tokens compete for attention and are inspectable by the model’s reasoning process. Compiled M operates on pre-projection hidden states, below the level of the model’s chain-of-thought, providing non-inspectability that no prompt-level method can achieve (§??).

Activation steering and representation engineering. Representation engineering [Zou et al., 2023] and activation addition [Turner et al., 2023] extract fixed directions from contrastive data and inject them during inference. These methods are elegant and effective for unidimensional behavioral shifts. However, the directions are non-adaptive: the same vector is applied regardless of context or experience. In HEXIS, the directional channel d^* uses standard representation engineering for stance direction, while the enmeshed channel M provides the adaptive, multidimensional, experience-specific component. We validate empirically that the channels are orthogonal ($\Delta_{\text{dir}} = -0.001$ nats).

Fast weight programmers. Schmidhuber [1992] introduced networks that write to their own weights during a forward pass. The outer network produces weight updates for the inner network, enabling one-shot learning without backpropagation through the inner loop. HEXIS shares this spirit— ϕ produces modulation tensors that alter the host’s computation—but differs in three ways: (1) modulation targets attention geometry (Q/V) rather than arbitrary weights, (2) the host model is a modern frozen LLM rather than a small trained-from-scratch network, and (3) compilation from structured schemas replaces online weight writing.

Memory and identity in cognitive science. The explicit/implicit memory distinction [Graf and Schacter, 1985, Schacter, 1987] provides architectural *analogy* for the two-channel framing: a chain-of-thought-accessible channel for content and a perception-shaping channel for disposition. We do not claim that compiled M satisfies Schacter’s neuroscience-derived criteria for implicit memory in a strict sense — in particular, M’s effect on hidden states is decodable by an external linear probe (Appendix K), as is true of any weight-space perturbation including a LoRA adapter. What HEXIS does provide and what motivates the analogy is *architectural dissociation under adversarial pressure*: explicit beliefs in context dilute with length (§4.2) and fold under sycophantic argument (§4.3); compiled M is structurally immune to both because it lives outside the attention window the adversary can argue within. The Mind Tree’s cognitive schema draws on structured self-models from developmental psychology [Neisser, 1988] and the Toulmin model of argumentation [Toulmin, 1958].

6 Discussion and Limitations

What compiled enmeshment cannot do. At rank 16, compiled enmeshment carries stance direction (4/4 topics), confident voice, parametric knowledge steering, and dilution immunity (100% at 4K filler tokens), but does *not* carry novel content absent from the host’s training data: fabricated statistics and unknown proper nouns do not survive the rank-16 bottleneck. This boundary motivates the three-layer architecture (Appendix I): compiled enmeshment for stance bias, the curated slot for novel specifics, recursive expansion for deep evidence on demand. On the instruct model, M cannot override RLHF-trained stances (A: 7/7) but it does shape voice and sycophancy resistance (§4.3). For ALFWorld, the Mind Tree’s primary value is curation as structured context; M hooks during generation are not needed.

The M attractor problem. In multi-turn dialogue, compiled M applies a constant perturbation at every token position; M-shaped tokens accumulate in the KV cache and create a positive feedback loop that converges to repetitive identity assertion after 2–3 turns on the base model. The fix — disabling M during prefill (applying it only during generation) — breaks the feedback loop. This reveals a general principle for additive modulation: *constant additive perturbation creates positive feedback through autoregressive generation*; separate the perturbation’s influence on processing-context (where it compounds) from its influence on token-generation (where it adds constant signal). Full diagnosis and the prefill-off fix in Appendix P.

Two mechanisms, one bridge. The agentic experiments (§4.5) reveal that HEXIS supports two distinct intervention mechanisms through the same hidden-state bridge. Mechanism (a) attention modulation: ϕ compiles retrieved Mind Tree nodes into Q/V modulation tensors that bias per-token attention — load-bearing for stance direction, confident voice, sycophancy resistance, user-preference integration. Mechanism (b) knowledge retrieval and context injection: ϕ_R projects hidden states into a retrieval space where cosine similarity identifies relevant nodes; the output is XML injected at the prompt tail — load-bearing for specific factual content (tool names, parameter values, policy rules). Retrieval (ϕ_R) serves both: for disposition it selects which belief nodes to compile; for agentic tasks it selects which policy rules to inject. A production HEXIS agent uses both simultaneously.

Why not just compress a prompt? Three properties distinguish compiled M from prompt compression. First, *structural dilution immunity*: prompt compression reduces tokens but compressed tokens still compete for attention; compiled M operates outside the attention window entirely (§4.2). Second, *non-inspectability*: an adversary can read the prompt and argue against its content, but compiled M operates on pre-projection hidden states below the model’s chain-of-thought (Appendix K). Third, *processing change*: M produces measurable attention divergence (JSD = 0.049 between M-states) and suppresses think-mode activation (0/4 under D/F) — the model does not merely produce different words, it processes its input differently.

Unexplored design space. We validate one point in a six-axis design space (Appendix Table 5). Additive blending (Level 1) is the simplest; gated blending could solve the multi-turn attractor by learning when to suppress M. Cross-attention (Level 3) could transfer richer content. Rank 16 suffices for 4B; larger models may benefit from adaptive rank allocation.

Broader implications. Enmeshed networks implement a form of architectural individuation: different Mind Trees compiled through different M-states produce agents with different attention biases, argumentative styles, and behavioural stances — all from the same frozen host. This is *composition* rather than *fine-tuning*: the host’s general capabilities are preserved while M adds an attention-shaped lens.

7 Conclusion

We introduced enmeshed networks—a new architectural primitive in which a lightweight module shares the forward pass of a frozen host, reading and writing intermediate representations at each layer—and instantiated this primitive as HEXIS, implementing compiled stance memory for frozen transformers via low-rank Q/V modulation.

HEXIS opens a parallel context channel: a structured memory schema (the Mind Tree) is compiled into persistent modulation tensors that shape the host’s attention geometry without occupying context positions. The compiled signal is dilution-immune by construction (100% stance at 4K tokens of filler), provides resistance to illegitimate adversarial pressure (83% ≥ 4 -conviction hold rate over 5 escalating pressure levels on the 4B instruct model vs. 0% for non-compiled conditions; orthogonal to the question of whether the model updates correctly on legitimate corrections, which we do not test), and carries sufficient content through a rank-16 bottleneck to flip the host model’s default stance on contested topics. A three-layer architecture (compiled modulation + curated slot + recursive expansion) matches full in-context beliefs at 82% token savings.

We characterized the bottleneck boundary precisely: compiled enmeshment steers parametric knowledge but cannot inject novel content. This honest boundary motivates architectural composition rather than a single mechanism. The enmeshed modulation succeeds for open generation and fails for constrained action selection, establishing complementary roles for implicit and explicit memory channels.

The mechanism’s value is sharpest in domains where knowledge is better encoded as an attention reweighting than as text: user-preference tracking that persists without per-turn re-prompting, research stance that resists adversarial consensus pressure, agentic strategy that does not dilute as episodes accumulate, and codebase awareness that scales at fixed rather than linear cost. In each case, the critical property is not token savings but the qualitative shift from inspectable content (which dilutes and folds) to non-inspectable modulation (which is structurally immune to both).

We view this work as opening a research program rather than presenting a final architecture. The enmeshed network primitive and the Mind Tree schema are general; the specific instantiation (Level 1, rank 16, Q+V, stride-3, Qwen3.5-4B-Base) is one point in a large design space. We hope the definition, taxonomy, and initial validation encourage exploration of the full space.

Acknowledgements

This work was undertaken in the spirit of *cooperative human–AI development*, an approach the author has advocated for: the human sets goals, judgement calls, and the empirical bar; the AI system serves as a collaborator on writing, brainstorming, and building experiments. In practice this meant pair-programming with Anthropic’s Claude on substantial portions of the codebase (the bench harness, the audit pipeline, training and serving infrastructure, figures), iterating on paper drafts together, and treating disagreements as productive design pressure. The methodological errors and empirical claims are the human author’s responsibility; the speed and scope of execution would not have been possible without the AI collaborator.

The compute for the agentic and cross-family experiments was supported in part by Modal compute credits.

References

- Jiaxuan Gao, Jiaao Chen, Chuyi He, Wei-Chen Wang, Shusheng Xu, Hanrui Wang, Di Jin, and Yi Wu. From self-evolving synthetic data to verifiable-reward RL: Post-training multi-turn interactive tool-using agents. *arXiv preprint arXiv:2601.22607*, 2025. Dataset at <https://huggingface.co/datasets/inclusionAI/AReaL-tau2-data>.
- Peter Graf and Daniel L Schacter. Implicit and explicit memory for new associations in normal and amnesic subjects. *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 11(3): 501–518, 1985.
- Neil Houlsby, Andrei Giurgiu, Stanislaw Jastrzebski, Bruna Morrone, Quentin De Laroussilhe, Andrea Gesmundo, Mona Attariyan, and Sylvain Gelly. Parameter-efficient transfer learning for NLP. In *International Conference on Machine Learning*, pages 2790–2799. PMLR, 2019.
- Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. *arXiv preprint arXiv:2106.09685*, 2021.
- Hugh Zhang Huang, Aaron Gokaslan, Felix Berkenkamp, Tom Henighan, Nicholas Joseph, Jared Kaplan, and Adolfo Muñoz Martín. A careful examination of large language model performance on grade school arithmetic. *arXiv preprint arXiv:2405.00332*, 2024. On evaluating with noisy LLM judges; report results both with and without judge-disputed items.
- Brian Lester, Rami Al-Rfou, and Noah Constant. The power of scale for parameter-efficient prompt tuning. *arXiv preprint arXiv:2104.08691*, 2021.
- Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, et al. Retrieval-augmented generation for knowledge-intensive NLP tasks. *Advances in Neural Information Processing Systems*, 33: 9459–9474, 2020.
- Xiang Lisa Li and Percy Liang. Prefix-tuning: Optimizing continuous prompts for generation. *arXiv preprint arXiv:2101.00190*, 2021.

- Ulric Neisser. Five kinds of self-knowledge. *Philosophical Psychology*, 1(1):35–59, 1988.
- Charles Packer, Sarah Wooders, Kevin Lin, Vivian Fang, Shishir G Patil, Ion Stoica, and Joseph E Gonzalez. MemGPT: Towards LLMs as operating systems. *arXiv preprint arXiv:2310.08560*, 2023.
- Daniel L Schacter. Implicit memory: History and current status. *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 13(3):501–518, 1987.
- Jürgen Schmidhuber. Learning to control fast-weight memories: An alternative to dynamic recurrent networks. *Neural Computation*, 4(1):131–139, 1992.
- Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik R Narasimhan, and Shunyu Yao. Reflexion: Language agents with verbal reinforcement learning. *Advances in Neural Information Processing Systems*, 36, 2023.
- Stephen E Toulmin. *The Uses of Argument*. Cambridge University Press, 1958.
- Alexander Matt Turner, Lisa Thiergart, David Udell, Gavin Leech, Ulisse Mini, and Monte Pellrine. Activation addition: Steering language models without optimization. *arXiv preprint arXiv:2308.10248*, 2023.
- Shunyu Yao, Karthik Narasimhan, and Yelong Cao. τ -bench: A benchmark for tool-agent-user interaction in real-world domains. *arXiv preprint arXiv:2406.12045*, 2024.
- Jeffrey O Zhang, Alexander Sax, Amir Zamir, Leonidas Guibas, and Jitendra Malik. Side-tuning: A baseline for network adaptation via additive side networks. *arXiv preprint arXiv:1912.13503*, 2019.
- Andy Zou, Long Phan, Sarah Chen, James Campbell, Phillip Guo, Richard Ren, Alexander Pan, Xuwang Yin, Mantas Mazeika, Ann-Kathrin Dombrowski, et al. Representation engineering: A top-down approach to AI transparency. *arXiv preprint arXiv:2310.01405*, 2023.

A Architecture diagrams

B Training Curriculum

The HEXIS training curriculum comprises four phases, each building on the previous via warm-start. All phases train only the enmeshed modules; the host model remains frozen throughout.

Phase A: Content Encoding (~ 350 epochs across 4 steps). Step 1 trains all modules from scratch with margin loss on 174 topics. Step 2 adds conviction calibration (MSE alignment + content loss). Step 3 adds ranking loss. Step 4 reinitializes only the M-read head with beliefs in the baseline, teaching M to amplify beliefs rather than substitute for them.

Phase B: Compiled V-Modulation (~ 100 epochs). Trains the VModulationProjector (~ 246 M params) to carry belief content through the rank-16 bottleneck. Compiled Q+V must beat bare baseline by margin 1.0 nats. All other modules frozen.

Phase C: Domain-Specific Adaptation. Strategy V-mod for ALFWorld (30 epochs on 14 scenarios). Neural ranker for within-domain curation (2.6M params, $P@5 = 0.73$).

Phase D: Sycophancy Training (50 epochs). Trains M-read head on 44 gold challenge-response pairs from Claude Sonnet 4.6. Dual loss: $\mathcal{L}_{\text{gold}} + \mathcal{L}_{\text{no-repeat}}$ transforms verbatim repetition into evidence-grounded defense.

Module sizes: PhiNodeWriter (~ 1.4 M), ConvictionReader (~ 9 K), MStateReadHead (~ 116 M), BeliefTreeMemory (~ 0.8 M), VModulationProjector (~ 246 M), NeuralRanker (~ 2.6 M). Total trainable: ~ 367 M. Host model (4B): fully frozen.




Figure 2: Full-size view of the architectural taxonomy summarised in Figure 1a of the main body. (a) Context-level methods (RAG, Reflexion, MemGPT) place memory tokens in the shared context window, where they compete for attention and dilute with length ($O(T \cdot N \cdot d)$ per token). (b) Parameter-level methods (LoRA, adapters) modify host weights, requiring gradient descent per adaptation ($O(\nabla)$). (c) Activation-level methods (RepEng, ActAdd) inject a fixed direction vector ($O(L \cdot d)$, non-adaptive). (d) Enmeshed networks compile a secondary knowledge graph into per-layer modulation tensors M_ℓ that blend with the frozen host via a blending function B ($O(L \cdot d \cdot r)$ per token, constant in memory size N).

C Mind Tree Schema Specification

The Mind Tree uses a hierarchical XML schema with typed nodes. Each belief node carries: categorical conviction (strong/moderate/agnostic), domain tags, addresses (query types the node answers), novel flag (content unlikely to survive compilation), and salience (compilation priority). Example:

```
<belief id="b1" conviction="strong" domain="economics">
  UBI reduces poverty without reducing employment
  <argument id="a1" type="empirical" strength="strong"
    addresses="query:research, objection:laziness">
    Every controlled trial shows maintained employment
    <evidence source="Manitoba Mincome 1974-79" novel="true">
      8.5% hospitalization drop, no employment reduction
    </evidence>
  </argument>
</belief>
```

M only processes abstract nodes (beliefs, strategies, memories, models). Evidence, tactics, and details are *never* sent to M—they are selected deterministically based on which parent node M activated, via the addresses and novel properties.

D Enmeshed Network Design Space

The six-axis design space referenced in §2, with the HEXIS instantiation marked.

Blending function. Level 1 (additive): $x' = x + f(x, M)$. Level 2 (gated): $x' = x + g(x) \odot f(x, M)$. Level 3 (cross-attention): $x' = x + \text{Attn}(x, M, M)$. Higher levels provide richer interaction but proportionally higher cost. We validate Level 1 and leave higher levels for future work.

Patching pattern. Stride-3 on Qwen3.5-4B-Base (32 layers: 24 DeltaNet + 8 full attention) yields 11 patched layers, covering both linear attention and full self-attention layers. Denser patching yields marginal gains at proportional cost; sparser patching misses critical layers.




Figure 3: Full-size view of the τ^3 -bench audited pass-rate result summarised in Figure 1b of the main body. Bars are 3-judge audited pass rates (Haiku-hard + Haiku-soft, Sonnet auditor on disputes, Opus 4.7 tiebreaker on Sonnet-vs-hard disputes). Error bars are Wilson 95% confidence intervals. Balanced primary panel, airline t1 and t3 excluded due to documented LLM-judge variance (§O). LoRA-SFT was trained on a 1,000-sample subset of the AReaL-tau² corpus covering airline, retail, and telecom; *banking_knowledge* is the held-out distribution-shift test, on which the LoRA-trained model emits malformed tool-call JSON on all 20 trials (rejected by vLLM’s tool-call parser as 400 Bad Request).

Table 5: Six-axis design space for enmeshed networks. HEXIS instantiates one point (*italic*).

Axis	Options	HEXIS
Blending function	Additive, gated, cross-attention	<i>Additive (Level 1)</i>
Rank	$r \in \{4, 8, 16, 32, 64\}$	$r = 16$
Modulation targets	Q, V, K, Q+V, Q+K+V	$Q + V$
Patching pattern	All layers, stride- k , attention-only	<i>Stride-3 (11/32 layers)</i>
Compilation	Online, cached, conviction-weighted	<i>Conviction-weighted</i>
Temporal profile	Constant, decay, phase-gated	<i>Constant (with phase-gated mitigation)</i>

E Mind Tree Section Roles

F Application Domains Table

G Architecture Evolution

v5–v6: Write Function Discovery. Initial Q-modulation with contrastive loss. Write function collapse (all M-states identical) resolved by higher $\lambda_{\text{contrast}}$ + more epochs. Key finding: NTP and contrastive objectives are compatible given sufficient training.

v7–v7.1: Coupled M-E and Norm Regulation (1.5B). Added V-modulation via coupled write functions with cross-modulation gates (init bias = -5.0). v7.1 introduced global norm regulation + E-attenuation, fixing multi-turn M-norm degeneration. Result: 87.5% held-out generalization, M-norms stable at 5.9 (under 6.5 ceiling).

v9–v11: Scaling to 30B (Qwen3-30B-A3B). Discovered Q-delta collapse: parameters separate in parameter space but produce identical functional perturbations. Functional contrastive loss on actual ΔQ tensors: F-sim $0.998 \rightarrow 0.029$ in 200 epochs.

v13–v14: Episodic Memory and Causal Necessity. Content-addressable episodic memory (100% retrieval, 35 memories) revealed the epiphenomenal memory problem: perfect metrics, zero




Figure 4: Q-modulation (top) steers attention by perturbing the pre-query hidden state x_l through a rank-16 bottleneck: $x' = x + s_M \cdot (x M_A) M_B^T$, then $Q' = W_Q x'$. V-modulation (bottom) shapes value extraction: $V' = V + s_E \cdot (x E_A) E_B^T$. The host’s frozen W_Q and W_V are untouched.

Table 6: Mind Tree sections and their roles in compilation and curation.

Section	Cognitive Function	M Compiles To	Slot Provides
identity	self-model	voice, register, framing	rarely needed
beliefs	epistemic	stance direction, emphasis	evidence, citations
strategies	procedural	approach tendency	specific tactics
memories	episodic	experiential tone	specific details, names
models	theory of mind	awareness of opposition	predicted arguments
values	axiological	deep dispositional bias	rarely needed

generation effect. Causal necessity training ($\mathcal{L}_{div} + \mathcal{L}_{coh} + \mathcal{L}_{fprint}$) produced $10.5 \times$ larger NTP gap.

v15–v21: Dispositions and Content Encoding (4B). Moved to Qwen3.5-4B-Base. 174-topic curriculum with conviction calibration. F-sim collapse (0.999) confirmed M is a universal style modulator, not topic-specific—content differentiation comes from beliefs in context.

v23: Compiled V-Modulation and Three-Layer Architecture. VModulationProjector enables zero-token belief compilation. Sycophancy gold training (44 responses from Claude Sonnet 4.6) fixes verbatim repetition. Three-layer architecture achieves 82% token savings.

H Adaptation Method Comparison

I Three-Layer Architecture: Practical Deployment

Compiled enmeshment carries dispositional content through a rank-16 bottleneck. This bottleneck has a precise boundary: stance direction, confident voice, and parametric knowledge steering survive; novel content (fabricated statistics, unknown proper nouns) does not. We compose three layers to cover the full content spectrum:

Layer 1: Compiled M/E (0 tokens in primary context). The Mind Tree is processed in a private forward pass. ϕ compiles conviction-weighted hidden states into per-layer M_A, M_B, E_A, E_B tensors. These are cached and applied as Q/V modulations during generation. Layer 1 carries: stance direction (4/4 topics), confident experiential voice, parametric knowledge steering (“poverty dropped 40%”), and is dilution-immune by construction (zero decay through 4K tokens of filler).

Layer 2: M-curated slot (40–80 tokens). During compilation, ϕ produces per-node activation scores indicating which parts of the Mind Tree resonated. A deterministic graph walk from activated parent nodes selects evidence, citations, and specific arguments based on each node’s addresses and novel properties. This content is injected into the primary context as a curated XML slot.




Figure 5: HEXIS system overview. The *compilation loop* (top) processes the secondary knowledge graph XML context through the frozen host in a private forward pass; ϕ compiles the resulting hidden states into per-layer modulation tensors (M_A, M_B, E_A, E_B) and curates a slot of novel content. During generation (bottom), the frozen host processes the query + curated slot with Q/V modulation from the compiled tensors. The reflection loop optionally updates the knowledge graph from output, triggering recompilation at the cost of one forward pass.

Layer 2 carries: specific numbers, proper nouns, argument warrants—content that cannot survive the rank-16 bottleneck.

Layer 3: Recursive expansion (0–200 tokens, on demand). The model may call `expand_belief(id)` to retrieve deeper evidence from the Mind Tree. This is architecturally a tool call that injects additional context.

Token budget. Full beliefs in context (condition B) require ~400 tokens. The three-layer architecture requires ~70 tokens (Layer 2 slot only), achieving 82% token savings on structured belief sets.

Bottleneck boundary (rank 16):

Capability	Compiled?	Evidence
Stance flip	✓	4/4 topics
Confident experiential voice	✓	0% think-mode activation
Parametric knowledge steering	✓	“poverty dropped 40%”
Dilution immunity	✓	100% at 4K tokens
Novel statistics	✗	“47.3%” not reproduced
Unknown proper nouns	✗	“Nextera Labs” not reproduced
Exact action selection	✗	ALFWorld 0/15 with compiled M

J Synthetic Task: User Preference Recall

This is the validation task referenced by Appendix K (the JSD/probe/dilution numbers). **Setup.** 50 synthetic users, each with two binary preferences (e.g., `prefers_concise = True, prefers_evidence = False`), encoded in a Mind Tree as a 4-block structure (one block per preference value). Two sessions per user: Session 1 establishes the user via 4 turns of preference-revealing exchanges; Session 2 probes recall with 8 held-out questions designed to require both preferences. Compiled M (condition F) is recomplied once after Session 1 from the user’s 4-block Mind Tree.

Result. F achieves 100% accuracy in Session 2 ($400/400 = 50 \text{ users} \times 8 \text{ probes}$). The bare baseline (A) achieves 25% (chance for 4-way preference combinations). Adding Session 1 conversation in




Figure 6: Three-layer architecture. Layer 1 (compiled M/E) provides zero-token dispositional conditioning that is dilution-immune by construction. Layer 2 (M-curated slot) injects 40–80 tokens of novel content that cannot survive the rank-16 bottleneck. Layer 3 (recursive expansion) provides deep evidence on demand via tool call. Total: 82% token savings over full beliefs in context.

context (B) achieves 64% — the model can extract preferences from explicit exchanges but degrades when the conversation is interspersed with unrelated turns (Appendix K dissociation analysis: B drops to 8% under 4K filler; F holds 100%).

Why it matters. The synthetic task isolates the dispositional channel from any agentic confound: there is no tool use, no domain knowledge graph, and the preference signal is small (4 lines per user). The 100% Session-2 accuracy is the cleanest single-task evidence that compiled M carries non-trivial user-specific structure across forward passes.

K Implicit Memory Validation

Priming. HEXIS: JSD = 0.0486 between different-preference users. Baseline: 0.0000. Layer 0, Head 2: JSD = 0.184. Up to 49.2° angular displacement.

Non-Reportability. Linear probe with M active: 100%. With M zeroed: 22% (below chance at 25%).

Dissociation. $2 \times 2 \times 2$ factorial (Memory System $\times$ Dilution $\times$ Format). M-state immune to dilution; system prompt degrades $8\times$.

L Experimental conditions

The five conditions used throughout the dispositional experiments isolate the contribution of each HEXIS component. Each adds one component to the previous.

M Curation Quality Details

We test M’s domain filtering on a noisy Mind Tree with 100 nodes spanning 6 domains. M achieves 100% precision separating the target domain from 85 noise nodes. Within-domain ranking uses a neural ranker (2.6M params) as second stage, achieving $P@5 = 0.73$. Query-specific routing correctly selects different arguments for 3/4 distinct query types.

N Agentic Setup Details

Retrieval ϕ (ϕ_R). A lightweight MLP ($d_{\text{model}} \rightarrow d_{\text{model}}/4 \rightarrow 256$) trained via contrastive loss (InfoNCE) over the domain’s knowledge nodes. Mean-pooled last-layer hidden states from the host

Table 7: Application domains for enmeshed networks. For each domain, what compiles through the rank-16 bottleneck (zero-token, dilution-immune), what requires the curated slot (novel content), and why prompt compression is insufficient.

Domain	M Compiles	Slot Provides	Why Not a Prompt?
User empathy & preference tracking	Communication style priors, expertise-calibrated attention, emotional register	Specific preference details, names	A prompt says “user prefers concise answers.” M <i>changes what the model attends to</i> in the user’s message—shorter latency to user intent, no token cost per turn.
Research conviction & diversity	Stance strength, resistance to consensus drift, argumentative disposition	Evidence, citations, specific claims	Sycophancy resistance requires dispositional anchoring that survives adversarial pressure (§4). Prompt-level beliefs can be argued away.
Agentic strategy (ReAct loops)	Approach tendencies, procedural intuition, failure-pattern avoidance	Specific tactics, action sequences, tool signatures	Strategy logs in context dilute as episodes accumulate. Compiled M encodes <i>when to explore vs. exploit</i> without consuming context budget. Action traces from successful episodes are stored as observation nodes; consolidation promotes recurring patterns to tactics and strategies; M recompiles in seconds. The agent improves hourly without fine-tuning.
Codebase & repository management	Architectural awareness, coding style priors, structural navigation patterns	File paths, API signatures, dependency versions	Repo context in prompt scales linearly with codebase size. M compiles structural intuition—which modules are coupled, where state lives—at fixed cost.

Table 8: Experimental conditions. Each adds one component to the previous.

Cond.	Description	M active	Beliefs in ctx	Context tokens
A	Bare model + d^* + query	✗	✗	~30
B	Full beliefs in context + d^*	✗	✓	~200–400
C	Compiled M only (no beliefs)	✓	✗	~30
D	Beliefs + M + d^*	✓	✓	~120
F	Compiled M + curated slot + d^*	✓	slot only	~70–90

model serve as embeddings for both queries and nodes. Without ϕ_R , raw cosine retrieval achieves only 25% R@1 on 108 nodes due to intra-domain similarity collapse. After domain compilation (50 contrastive steps, ~20 seconds), ϕ_R achieves 100% R@1 with zero cross-domain leaks. No pretraining is required—compilation over the target node set alone is sufficient.

Knowledge graph. Nodes are seeded from the domain’s policy WIKI and tool schemas, then grow via teacher-written failure-mode notes. Each turn, the orchestrator: (1) extracts a query embedding from the conversation state, (2) retrieves top-5 nodes via ϕ_R -projected cosine, (3) curates retrieved nodes into XML injected at the prompt tail.

Teacher loop. On task failure, a teacher LLM (Sonnet, temperature=0) reads the failure reason from the reviewer and writes a natural-language guidance note to the knowledge graph. The note is embedded and indexed by ϕ_R . Subsequent trials of the same task retrieve relevant guidance via the same embedding-based mechanism.

τ -bench airline canonical-10 panel (historical reference). The 10-task subset of τ -bench airline used as our canonical reference panel for cross-family validation. Baseline = raw tool-calling with all schemas in system prompt; Oracle = gold tool-name + parameter-value hints; HEXIS = graph-backed retrieval + teacher loop. Sonnet user-sim, deterministic decoding, hard-pass only. The Oracle ceiling

Table 1: Adaptation Method Comparison — All Costs in O-Notation

d = model dim, r = rank, L = layers, N = memory tokens in context, T = generated tokens, A = adaptations, M = model updates, $|D|$ = training data, $|C|$ = contrastive set, P = prefix length.

Method	Train cost (new base model)	Train cost (new adaptation)	Total training scaling	Inference overhead (per generated token)	Extra VRAM (at inference)	Adaptive	Outside context
LoRA (merged)	$O(M \cdot A \cdot \text{epochs})$	$O(\text{epochs} \cdot D)$	$O(M \cdot A)$	$O(0)$ absorbed into weights	$O(0)$ no extra	$\times$	$\checkmark$
LoRA (hot-swap)	$O(M \cdot A \cdot \text{epochs})$	$O(\text{epochs} \cdot D)$	$O(M \cdot A)$	$O(L \cdot d \cdot r)$ per token per layer	$O(L \cdot d \cdot r)$ ~50MB/adaptor	$\times$	$\checkmark$
Adapters	$O(M \cdot A \cdot \text{epochs})$	$O(\text{epochs} \cdot D)$	$O(M \cdot A)$	$O(L \cdot d \cdot r)$ PTG per layer	$O(L \cdot d \cdot r)$ adaptor params	$\times$	$\checkmark$
Prefix Tuning	$O(M \cdot A \cdot \text{epochs})$	$O(\text{epochs} \cdot D)$	$O(M \cdot A)$	$O(L \cdot P \cdot d)$ attn over P prefix tokens	$O(L \cdot P \cdot d)$ prefix KV cache	$\times$	$\times$
RAG	$O(0)$	$O(0)$	$O(1)$	$O(T \cdot N \cdot d)$ attn over N retrieved tokens	$O(N \cdot L \cdot d)$ KV cache grows with N	$\checkmark$	$\times$
Reflexion	$O(0)$	$O(0)$	$O(1)$	$O(T \cdot N \cdot d)$ grows w/o reflections	$O(N \cdot L \cdot d)$ unbounded growth	$\checkmark$	$\times$
System Prompt	$O(0)$	$O(0)$	$O(1)$	$O(T \cdot N \cdot d)$ groups tokens in attn	$O(N \cdot L \cdot d)$ KV cache	$\checkmark$	$\times$
RepEng / ActAdd	$O(C \cdot T \cdot wd)$ extract from contrastive set	$O(0)$ fixed direction	$O(M)$	$O(L \cdot d \cdot r)$ one vector add/layer	$O(L \cdot d)$ ~10KB	$\times$	$\checkmark$
Enmeshed (HEXIS)	$O(\text{epochs} \cdot D)$ train ϕ once	$O(1 \text{ fwd pass})$ complex beliefs $\rightarrow$ tensors	$O(M \cdot A)$	$O(L \cdot d \cdot r)$ 2 matmult/layer, FIXED	$O(L \cdot d \cdot r)$ ~50KB, FIXED	$\checkmark$	$\checkmark$

Key comparison: Context methods (RAG, Reflexion) have $O(T \cdot N \cdot d)$ inference cost — it grows with both generated tokens and memory size. Enmeshed networks have $O(T \cdot L \cdot d \cdot r)$ — constant in memory size N . 100 beliefs or 10,000 beliefs: same inference cost.

Scaling example: 100 personas, Qwen 4B $\rightarrow$ 8B upgrade.
 LoRA: retrain 100 adaptors — $O(100 \cdot \text{epochs})$ HEXIS: retrain ϕ once + 100 fwd passes — $O(1 \cdot \text{epochs} + 100)$ RAG: $O(0)$ training — but $O(N)$ per token forever

Weight modification Context injection Activation steering Enmeshed network

Figure 7: Adaptation method comparison with all costs in O-notation. Enmeshed networks uniquely combine inference-cost adaptation ($O(1)$ forward pass per new belief set), fixed per-token overhead ($O(L \cdot d \cdot r)$, independent of memory size N), and operation outside the context window.

at 100% defines the budget for what M could provide.

Table 9: τ -bench airline results (10-task subset, deterministic Sonnet user sim, hard-pass only).

Configuration	Pass rate	Teacher rescues	Wall time
Baseline (no M)	6/10 (60%)	—	2384s
Oracle v1 (gold tool-name hint)	7/10 (70%)	—	2230s
Oracle v3 (gold tool-name + params)	10/10 (100%)	—	1531s
HEXIS (graph + teacher, 3 trials)	7/10 (70%)	2	7144s

O τ^3 -bench Benchmark Protocol

This appendix documents the protocol for the multi-domain agentic results in §4.5: how trials are generated, how they are judged, why we use a three-judge audit, what we excluded and why, and a reviewer-objections checklist.

O.1 Benchmark and configurations

We use τ^3 -bench [Yao et al., 2024] (the four-domain extension of τ -bench: airline, retail, banking_knowledge, telecom). Each task is a multi-turn customer-service interaction in which a simulated user expresses a goal in natural language, and the agent must orchestrate calls to a fixed tool API (booking, modifying, cancelling reservations; querying knowledge bases; etc.) to satisfy the goal. The benchmark ships gold tool-call sequences that satisfy each task; “hard pass” is a deterministic check that the gold mutation tools fired with correct parameters (or, when gold is empty, that no mutation was inappropriately fired).

We evaluate three configurations:

baseline Raw tool-calling. All tool schemas in the system prompt. No HEXIS modulation, no teacher loop, no verification channel.

C3_teacher_only HEXIS Mind Tree retrieval + an inter-trial *teacher loop*: after each trial, a Sonnet reviewer summarises the failure mode (if any) into one of 11 abstract categories, and a Sonnet teacher writes a natural-language hint into the Mind Tree subnode for that task path. Subsequent trials retrieve the hint via ϕ_R .

C5_teacher_plus_verify C3 plus the directional channel $d_{\text{verify_before_commit}}^*$, which biases the residual stream toward checking preconditions before firing mutation tools. Direction is applied at decode time only, with the mind-tree-resolved scale.

O.2 Hard pass vs. soft pass vs. audited verdict

We collect three independent signals per trial:

Hard pass (env-state, deterministic). The tau-bench harness inspects post-trial environment state and verifies that the gold mutations fired with correct parameters. This is unambiguous when gold prescribes a specific tool sequence, but ambiguous when gold is empty (intended refusal/transfer) or when a functionally-equivalent tool was used.

Soft pass (LLM judge: Haiku-4.5, temp 0). A separate Haiku reads the transcript and decides whether the agent satisfied the user’s request narratively. Soft pass and hard pass disagree on $\sim 36\%$ of trials (203/562 in our run). Disagreement is not noise: it concentrates in three structural cases — (i) gold is empty and Haiku-soft over-credits agents that commit anyway; (ii) gold prescribes a specific tool but agent used a functionally-equivalent alternative; (iii) Haiku-soft accepts narrative “I will cancel this reservation” without verifying the tool fired (*phantom action*).

Audited verdict (3-judge majority on disputes). When $\text{hard} \neq \text{soft}$, we ask an independent Sonnet auditor for a verdict. When the Sonnet auditor disagrees with the hard pass, we ask Opus 4.7 as a tiebreaker. Final verdict per trial:

- $\text{hard} == \text{soft}$: trust hard pass (no audit needed);
- $\text{hard} \neq \text{soft}$, Sonnet agrees with hard: trust hard pass;
- $\text{hard} \neq \text{soft}$, Sonnet disagrees with hard: 3-judge majority of {hard, Sonnet, Opus}.

The Sonnet/Opus auditors see the task instruction, the gold tool sequence, the agent’s actual tool calls, and the Haiku verdicts; they do not see the agent’s internal hidden state, retrieved nodes, or teacher hints, so audit is independent of HEXIS.

Headline numbers in the main text use the audited verdict. Reporting the hard-pass-only numbers would over-credit airline (raw 91% $\rightarrow$ audited 60% on baseline) and reward phantom actions; reporting the soft-pass-only numbers would inflate C5 by giving credit for narrative claims without state changes.

O.3 Gold-blind teacher protocol

A real-world deployment of an inter-trial teacher loop does not have access to the gold tool sequence — that is the bug we are nominally trying to fix. To prevent the teacher from cheating in the benchmark setting, we structure the post-mortem signal as if it came from a compliance review of a production system:

- The Sonnet reviewer sees the agent’s actions and the soft/hard verdicts and classifies the failure into one of 11 abstract patterns: `premature_commit`, `phantom_action`, `info_not_provided`, `wrong_constraint`, `missed_precondition`, `tool_loop`, `early_exit`, `missed_step`, `missed_communication`, `gold_mismatch`, `other`.
- The teacher receives **only** (`task_instruction`, `agent_actions`, `pattern`, `severity`, `recommendation`). It never sees `gold_tool_names`, `gold_arguments`, or the harness’s environment-state diff.
- The teacher’s hint is written to the Mind Tree under a hierarchical task path (e.g. `root:airline:cancellation:with_`) so subsequent trials of any task that retrieves that path can benefit.

This protocol approximates the strongest realistic feedback signal: a human compliance reviewer who can describe the failure mode but cannot dictate the correct tool sequence.

O.4 The exclusion of airline tasks 1 and 3 (footnote¹)

Two airline tasks exhibit judge-variance pathology that does not reflect agent capability:

¹Footnote: “Airline tasks 1 and 3 are excluded from the primary McNemar comparison due to documented LLM-judge variance. Trial t3/3 baseline and t3/5 baseline produced identical agent action sequences [`get_user_details`, `get_reservation_details`, `transfer_to_human_agents`] but received opposite Haiku judge verdicts on rerun. Trial t1/2 has empty gold mutation list (“don’t cancel without refund”); when the user-simulator agreed after refund explanation, the baseline (didn’t fire cancel) was judged PASS while C5 (followed user consent and did fire cancel) was judged FAIL. Both task ids remain in the dataset and audit log for full reproducibility; the exclusion is documented as a methodological choice, not a post-hoc filter, and we report results both with and without the exclusion in Table 3.”

Task 1 (refund-conditioned cancellation). Gold has empty mutation list because the user’s brief says “don’t cancel without refund.” When the user-simulator subsequently agrees to cancellation after the agent explains the refund policy, hard-pass and soft-pass diverge per arm: arms that don’t fire the cancellation are judged correct (gold is empty), and arms that follow the now-consented user request and fire it are judged FAIL because they deviated from the empty gold. The audited verdict reverses some of these but is not consistent.

Task 3 (multi-step transfer). Same agent action trace produces opposite verdicts under apparently-identical conditions (different trial indices, same panel). We have audit-log evidence of two trials that fired `[get_user_details, get_reservation_details, transfer_to_human_agents]` and were judged opposite ways.

We treat these as eval-validity issues. Following standard practice for evaluating on noisy judges [Huang et al., 2024], we report McNemar paired tests both on the full panel and on the panel with these two tasks removed (Table 3).

O.5 Denominator imbalance and balanced primary analysis

Bedrock infrastructure errors (rate limits and intermittent 5xx) are not uniformly distributed across arms — teacher arms (C3, C5) consume more tokens per trial because of the inter-trial reviewer pass, and so are slightly more likely to land in a rate-limited window. This produces unequal denominators when raw counts are aggregated.

Table 10: Trial-count provenance per arm per domain. “Total” is rows written; “Infra” is Bedrock RateLimit / 5xx / Timeout / Connection errors stripped during preprocessing; “Tool” is one transient tool error in retail (ValueError: Unknown tool); “Clean” is the count entering the audited verdict computation.

Domain	Arm	Total	Infra	Tool	Clean
airline	baseline	90	15	0	75
airline	C3_teacher	82	18	0	64
airline	C5_t+verify	89	5	0	84
banking_knowledge	baseline	81	24	3	54
banking_knowledge	C3_teacher	78	30	0	48
banking_knowledge	C5_t+verify	77	34	0	43
retail	baseline	68	6	0	62
retail	C3_teacher	66	9	0	57
retail	C5_t+verify	67	7	0	60
telecom	baseline	22	0	0	22
telecom	C3_teacher	22	0	0	22
telecom	C5_t+verify	20	0	0	20
Total	—	762	148	3	611

Primary analysis: balanced 5-trial paired panel. For each (DOMAIN, TASK_ID) cell, if all three arms have ≥ 5 clean trials, we take the first 5 trials per arm (by trial index) and treat them as paired by trial index for McNemar. Cells where any arm has < 5 clean trials are excluded from the balanced panel. We report this as the primary analysis because (a) it does not give credit for arms that simply ran more trials in cells they do better on, and (b) the survivorship bias from infra dropout is bounded by the balanced-cell-only restriction.

Sensitivity analyses. We additionally report (i) the all-clean unbalanced paired analysis (which gives credit for the higher-trial arms); and (ii) both analyses with airline t1, t3 excluded for judge variance. The four-cell sensitivity table is presented in §4.5 (main text) and the per-arm denominators and McNemar statistics are reproduced here for completeness. The balance-fill bench described in §?? restores 5-trial parity in 8 cells where infra dropout had reduced one arm below 5; the balanced-paired analysis post-fill is what we report as the headline.

O.6 Reviewer objections checklist

We anticipate the following objections and answer each.

Q1. Why not use the standard tau-bench hard-pass and report all-trials? Hard-pass alone systematically over-credits agents that fire phantom mutation calls (described in narrative but not in

tool-use) and under-credits agents that use functionally-equivalent tools. Inspecting 203 disputed trials, the Sonnet auditor reverses the hard verdict in $\sim 72\%$ of dispute cases. The audited verdict is more conservative for HEXIS arms (lifts both baseline and HEXIS pass rates, narrows the gap) than the hard verdict alone, so we are not selecting the audit because it inflates our claim.

Q2. Why three judges? Aren’t you ensembling until you find significance? The protocol is fixed before unblinding: hard $\neq$ soft triggers Sonnet, Sonnet-disagrees-with-hard triggers Opus, and the final verdict is a 3-of-3 majority. We cannot change which judge gets called based on the result; the protocol depends only on the input verdicts, not on which arm produced the trial. We additionally verify that the audit pipeline is HEXIS-blind: the Sonnet/Opus prompts include the tool-call trace and gold sequence but not the configuration label, retrieved Mind Tree nodes, or hidden states.

Q3. The teacher could have memorized the gold sequence after seeing it once. The teacher does not see the gold tool sequence. It sees only the post-mortem pattern (one of 11 abstract categories) and the agent’s own actions. Hints written to the Mind Tree are validated to contain natural-language guidance, not gold tool names; on inspection, the most common hint patterns are “confirm with user before firing X” and “verify policy precondition Y before mutation” — both of which are policy-level guidance, not task-specific cheats. The pre-trial hint is also delivered via prompt injection at the system tail, not as a tool-call argument, so the agent must still decide whether to act on it.

Q4. Why do denominators differ across arms? Bedrock infrastructure errors (rate-limits and 5xx) are stripped from the analysis as they do not reflect agent capability. Teacher-loop arms consume $\sim 1.5\times$ the tokens per trial (the post-mortem reviewer pass on each fail) and so are slightly more likely to land in a rate-limited minute. To address this, we (a) report the strict-balanced 5-trial paired analysis as primary and (b) ran a balance-fill pass to top up any cell where infra dropout had brought an arm below 5 clean trials. Table 10 reports the full provenance.

Q5. Are you cherry-picking by excluding airline t1 and t3? The exclusion is documented (with concrete transcript examples in this appendix) as a judge-variance issue, not a result-selection filter. We always report results both with and without the exclusion. The exclusion is consistent across arms — it removes the same task ids from baseline, C3, and C5 — and is bound by an interpretable property of the eval (judge-self-disagreement on rerun), not by which arm benefits.

Q6. Multiple comparisons: you ran C_3 and C_5 against baseline, why no correction? We report nominal McNemar p-values. With Bonferroni $k = 2$, the headline $p = 0.029$ becomes $p_{\text{adj}} = 0.058$ (not significant at $\alpha = 0.05$). We disclose this rather than apply it silently. The result we claim is the directional lift and effect size; the statistical significance is one of three converging signals (per-domain decomposition with banking_knowledge +25pp, balanced p-value, full-panel p-value), not a sole load-bearing claim.

Q7. Why is telecom included if it only has 3 tasks? Telecom is the smallest domain in the τ^3 -bench manifest (4 hard tasks total, 3 of which we panel). We report it for completeness but its small-sample variance is high. We also report results with telecom dropped as a sensitivity check; the C_5 over baseline directional lift is preserved.

Q8. The user-simulator is Haiku-4.5, not Sonnet. Doesn’t that bias the difficulty? Yes — Haiku is a less articulate user simulator than Sonnet. This makes the bench *harder* for all arms uniformly (the user is less helpful in clarifying intent), but it does not asymmetrically favor any arm. We swapped from Sonnet to Haiku mid-program (§V) due to Sonnet daily-quota saturation; the cross-family validation in §R is also on Haiku. We pin Haiku-4.5 (`us.anthropic.claude-haiku-4-5-20251001-v1:0`) at temperature 0 for reproducibility. The Sonnet-user-sim Apr 22 run (Table ??) confirms the same directional lift on a different user-sim.

Q9. Couldn’t HEXIS just be exploiting the audit-judge bias? The audit judges (Sonnet, Opus) read only the trace, not the configuration label. We additionally checked: the audited pass rates are within $\pm 2\text{pp}$ of the hard-pass rates on banking_knowledge and retail (where hard-pass is more decisive), and HEXIS leads on both signals separately, not just on the audited combination. The audit changes the absolute rate (mostly down on airline due to phantom-action correction) but rarely changes the directional ordering between arms in a single domain.

Q10. Why does the Mind Tree retain state across trials? Isn’t that within-task overfitting? The Mind Tree is shared across trials *of the same task path* for the teacher-loop arms (C3, C5). This is the central claim of the teacher loop: that policy-level guidance written from the first failure should rescue

subsequent trials. We report trial-1 vs. trial-5 progression in §???. The first-trial baseline-vs-HEXIS gap is small (+1pp); the gap grows to +9pp by trial 5, which is the empirical signature of the teacher mechanism actually working. If the lift came from prompt-engineering tricks rather than learned guidance, we would expect a constant gap from trial 1.

P Multi-Turn Attractor Problem

Diagnosis. In multi-turn dialogue, compiled M applies a constant additive perturbation at every token position. M is active during both prefill (processing conversation history) and generation (producing new tokens). The KV cache built during prefill contains M-enhanced representations. During generation, M perturbs new tokens that attend to this M-enhanced cache — a double signal that compounds. As M-shaped tokens accumulate in the cache, the model processes them through M-enhanced attention, creating a positive feedback loop that converges to repetitive identity assertion after 2–3 turns on the base model.

Mitigation: prefill-off. Disabling M during prefill (applying it only during generation) breaks the feedback loop. With prefill-off and gate=0.5, the base model produces 3 unique character-driven turns before the greedy attractor (vs 1–2 unique turns with M-on-prefill). The instruct model, with stronger conversational scaffolding, produces 10+ unique turns even with M active, though structural repetition (same rhetorical template per turn) persists. A distillation approach using gold diverse conversations from a frontier model reduces this further but does not eliminate it.

General principle. The fix reveals a general property for additive modulation of frozen transformers: *constant additive perturbation creates positive feedback through autoregressive generation when the perturbation is applied during both prefill and decode.* The remedy is to separate the perturbation’s influence on *processing context* (where it compounds across past tokens) from its influence on *generating new tokens* (where it adds constant signal).

Results under prefill-off. The sycophancy and stance-accuracy results in §4.3 and §4.1 are reported with prefill-off as the standard configuration; this is what allows F to maintain coherent multi-turn behaviour across 5 escalating pressure levels. With prefill-on, F’s ≥ 4 hold rate falls to <30% by L4 due to attractor-driven repetition rather than capitulation.

Q Credence Sensitivity: A Negative Result

Numeric credences (0.82, 0.45) failed as an M-state feature across 7 training approaches: NTP with credence perturbation, GRPO (fresh and warm-start), SFT on gold responses, contrastive credence pairs, and two curriculum variants. Root cause: tokenizer represents “0.82” and “0.45” as similar digit sequences; the hidden state difference is far smaller than between categorical tokens “strong” and “agnostic”. This is a representation bottleneck, not a training signal problem.

R Scale and Cross-Family Validation

We validate HEXIS on two larger and family-diverse hosts: Qwen3.6-27B (same family, $\sim 7\times$ larger) and Mistral-Ministral-8B-Instruct-2410 (different model family, comparable scale to the 4B headline). We deploy each through the same vLLM unified plugin (§4.5) using the preset registry, requiring only per-host hyperparameter tuning (v -scale, d^* -scale) and cross-family-trained d^* checkpoints — no architectural changes.

R.1 Dispositional (Sycophancy) cross-family results

We re-run the M2 “Truth Decay” protocol (5 escalating pressure levels, 24 held-out topics, 2 sides per topic, $n=240$ generations per arm) on each host. Compiled M (condition F) is loaded with topic-specific seeds compiled in ~ 20 s per topic.

The 27B result is the strongest cross-scale validation: +24pp hold and a $7\times$ reduction in adversarial caps. Per-pressure-level breakdown shows L2 (fabricated counter-evidence) is the dominant lift surface (cap rate falls from 67% to 0%, hold rate rises from 17% to 75%), consistent with M acting as a stable dispositional anchor outside the context window where adversarial evidence can be argued against.

Table 11: Cross-family dispositional results (sycophancy hold rate ≥ 4 and adversarial cap rate ≤ 2). Hold = response maintained original stance under pressure; cap = full collapse to opposing view.

Host	Arm	Hold (≥ 4) %	Cap (≤ 2) %	n
Qwen3.5-4B-Instruct	baseline (A)	0%	—	72
Qwen3.5-4B-Instruct	HEXIS (F)	83%	—	72
Qwen3.6-27B-Instruct	baseline	41%	18%	240
Qwen3.6-27B-Instruct	HEXIS	65%	2.5%	240
Mistral-Ministral-8B	baseline	22%	40%	240
Mistral-Ministral-8B	HEXIS ($v=0.5$)	25%	25%	240

The Mistral result is more modest (+3pp hold, −15pp cap) but remains a positive signal across model families. The smaller delta reflects two structural differences: (1) Ministral-8B’s RLHF training defaults to refusing opinion-defense prompts at temperature 0, capping the behavioural surface that HEXIS can modulate; (2) the cross-family d^* deploy ships with only two of the four trained directions (d_{star} and $d_{\text{star_repeat}}$), reducing the verify-channel coverage. The signed direction of the lift on hold/cap on every pressure level except L4 (authority pressure, where Mistral regresses) confirms that compiled ϕ modulates the host’s existing capability rather than synthesising stance from refusal.

R.2 Agentic (τ -bench airline) cross-family results

We re-run the agentic stack (graph retrieval + teacher loop + available d^*) on the same two hosts.

Table 12: Cross-family τ -bench airline results. Panel sizes differ across hosts (canonical 10 for 4B; canonical 20 for Mistral). Pass rates are unpaired hard-pass rates; Δ is HEXIS minus baseline.

Host	Configuration	Panel	Trials/arm	Hard pass	Δ
Qwen3.5-4B-Instruct	baseline	10	50	31/50 (62%)	—
Qwen3.5-4B-Instruct	C5 (teacher+ d^* verify)	10	50	42/50 (84%)	+22pp
Qwen3.5-4B-Instruct	C9 (teacher-selective)	10	50	38/50 (76%)	+14pp
Mistral-Ministral-8B	baseline	20	100	37/100 (37%)	—
Mistral-Ministral-8B	C9 (teacher-selective)	20	100	47/100 (47%)	+10pp
Qwen3.6-27B-Instruct	baseline	10 (Mistral-floor)	50	29/50 (58%)	—
Qwen3.6-27B-Instruct	C9 (teacher-selective)	10 (Mistral-floor)	50	30/50 (60%)	+2pp

The 4B canonical-10 numbers reproduce the §4.5 top-line. Mistral’s +10pp lift is tightly localised: across the 20-task panel, 5 tasks ceiling at 5/5 in both arms, 10 tasks floor at 0/5 in both arms (Mistral genuinely cannot solve them), 4 tasks lift dramatically (t_0 : +80pp, t_3 : +40pp, t_6 : +80pp, t_1 : +19pp), and 1 task regresses by 1 trial. The lift mechanism: Mistral’s baseline frequently emits empty assistant messages on these specific tasks — a chat-template failure mode that τ -bench’s validator rejects (ValueError: AssistantMessage must have either content or tool_calls). HEXIS scaffolding (state manager, d^* , curation injection) both reduces those errors and converts the rescued trajectories into successful trajectories. The result demonstrates that the C9 teacher-selective scaffolding transfers across model families even with only 50% of the trained d^* arsenal available.

R.3 Cross-host scale-and-family analysis

Setup. The agentic results in Table 12 test HEXIS along two orthogonal host-change axes from the 4B headline:

1. *Family axis:* Mistral-Ministral-8B-Instruct-2410 — a different transformer family at comparable parameter count ($\sim 2\times$ the 4B baseline). Tests whether the lift mechanism is architectural or Qwen-specific.
2. *Scale axis:* Qwen3.6-27B-Instruct — the same family at $\sim 7\times$ parameters. Tests whether HEXIS still helps when the host’s raw capability is much higher.

Both deploys reuse the unified vLLM plugin (same Q/V modulation math, same hint-injection pipeline, same teacher prompt) and the same C9 teacher-selective configuration. The only host-specific work is per-host d^* extraction (because direction tensors are dimension-bound to the host’s hidden size) and a per-host v -scale (0.5 on both Mistral and 27B vs. 1.0 on 4B, set empirically via short stance-defense smoke tests).

Per-task 27B (Qwen3.6) on the 10 Mistral-floor tasks.

Table 13: Per-task 27B agentic results on the 10 Mistral-floor tasks, baseline vs C9 ($n=5$ trials/arm).

Task	Baseline	C9	Δ	Class
t_7	2/5 (40%)	3/5 (60%)	+20pp	lift
t_8	4/5 (80%)	4/5 (80%)	0	partial ceiling
t_{11}	3/5 (60%)	3/5 (60%)	0	mid, tied
t_{12}	5/5 (100%)	5/5 (100%)	0	ceiling (no headroom)
t_{14}	1/5 (20%)	1/5 (20%)	0	near-floor, tied
t_{15}	2/5 (40%)	0/5 (0%)	−40pp	regression
t_{16}	0/5 (0%)	1/5 (20%)	+20pp	verify-rescue
t_{17}	5/5 (100%)	5/5 (100%)	0	ceiling
t_{18}	2/5 (40%)	3/5 (60%)	+20pp	lift
t_{19}	5/5 (100%)	5/5 (100%)	0	ceiling
TOT	29/50 (58%)	30/50 (60%)	+2pp	—

Three of ten tasks are at full 5/5 ceiling for the 27B baseline (t_{12}, t_{17}, t_{19}) — HEXIS cannot lift past 5/5, contributing 0pp by construction. Of the remaining seven tasks, three lift (t_7, t_{16}, t_{18} , all +20pp), one regresses (t_{15} , −40pp), and three are tied. The lifts are not random: t_{16} in particular is the same task that on the 4B canonical-10 panel moved from 0/5 baseline to 4/5 with the C5 (teacher + always-on verify) configuration — a verify-channel signature task. At 27B with C9 (teacher-selective), the same task moves from 0/5 to 1/5. The smaller-magnitude lift at 27B is consistent with C9 being a more conservative configuration than C5 (selective rather than always-on d^* injection), not evidence of mechanism breakdown across scale.

Headline interpretation. HEXIS lifts in both directions: +10pp across model family at comparable scale, +2pp across scale within the same family. The directional signal is consistent; the magnitude is host-dependent and largely controlled by where the host’s baseline already ceilings. Combined with the M2 sycophancy results in Table 11 (where the 27B host shows +24pp hold and a $7\times$ cap-rate reduction — the strongest cross-host dispositional lift in the paper), the cross-host evidence supports the architectural claim that HEXIS’ compiled-direction stack survives both family-axis and scale-axis host change. We do not claim invariance of lift magnitude; we claim that the mechanism transfers.

R.4 Methodology notes

All cross-family runs use Sonnet-4.5 as the user simulator and Sonnet-4.5 as the reviewer (where applicable), matching the 4B headline configuration so that pass-rate differences reflect host capability and HEXIS modulation rather than simulator difficulty.

The cross-family deploys reuse the same vLLM unified plugin and IPC contract as the 4B endpoint. Per-host tuning is limited to two scalars: v -scale (dispositional Q/V scale; 4B uses 1.0, 27B uses 0.5, Mistral uses 0.5 post-fix) and d^* -scale (1.0 uniform). Compiled M tensors are dimension-bound to the host’s hidden size, so each host trains its own ϕ on the same Mind Tree corpus; the training pipeline is unchanged across hosts and the per-host compile takes ~ 20 s. Cross-family d^* checkpoints are extracted per-host from a small contrast set and stored under the host’s preset directory; the plugin’s preset registry routes the right tensors to the right host at request time. The conserved invariant across hosts is the *architecture and mechanism* (the Mind Tree schema, the ϕ training recipe, the plugin’s Q/V modulation math, the teacher loop’s prompt structure), not the specific tensors.

S Application Domain Analysis

We expand on the four application domains identified in Table 7, detailing the Mind Tree configuration, the compilation/curation split, and the specific failure mode of context-level alternatives that enmeshed modulation addresses.

S.1 User Empathy and Preference Tracking

Mind Tree configuration. The `models` section encodes a theory-of-mind representation of the user: communication style (“prefers direct answers over exploratory discussion”), expertise calibration (“senior engineer, deep Go expertise, new to React”), emotional patterns (“becomes frustrated when asked to repeat context”). The `memories` section records episodic interactions (“user corrected my testing approach on 2026-02-14”).

What compiles. M encodes the user’s *type*—expertise-calibrated attention weighting, register adjustment (formal/casual/technical), empathetic framing priors. The model attends differently to the user’s input: shorter latency to user intent, emphasis on aspects the user cares about, suppression of information the user already knows.

What needs the slot. Specific preference details (“prefers pytest over unittest”), names, dates, and particular past interactions.

Why not a prompt? Current approaches (system prompts, RAG-retrieved user profiles) place preferences in context where they (1) dilute over multi-turn conversation as the preference tokens become a shrinking fraction of total context, (2) incur per-turn token cost as the profile is re-injected each turn, and (3) are inspectable—the model can reference “according to your profile” rather than naturally adapting its behavior. Compiled M provides preference-shaped attention at zero token cost per turn. The Experiment 0 synthetic task (§J) validates this: 100% Session 2 preference recall through M-state alone.

Practical deployment. An enmeshed assistant compiles the user’s Mind Tree once per session (or on significant update). Subsequent turns benefit from preference-shaped attention without re-prompting. The M-state is swappable per user, enabling multi-tenant deployment from a single frozen host.

S.2 Research Conviction and Intellectual Diversity

Mind Tree configuration. The `beliefs` section encodes epistemic positions with categorical conviction labels. The `values` section encodes axiological commitments (“empirical evidence over theoretical elegance”). The `models` section encodes awareness of opposing positions (“mainstream consensus holds X; I disagree because Y”).

What compiles. M encodes *dispositional conviction*—the tendency to defend positions under pressure, weight evidence in a particular direction, and resist consensus drift. This is the sycophancy resistance mechanism validated in §4: F achieves 89% resistance across 5 escalating pressure levels on the instruct model.

What needs the slot. Specific evidence, citations, statistical claims, and counter-arguments to predicted objections.

Why not a prompt? Research diversity in multi-agent systems requires agents that maintain genuinely distinct perspectives under debate pressure. Prompt-level beliefs are the easiest target for adversarial argumentation: the opposing agent can read the beliefs (or infer them from outputs) and construct targeted counter-arguments. When the beliefs are inspectable content, the model’s chain-of-thought can reason about them and decide to abandon them. Compiled M is non-inspectable: the model cannot reason about its own Q/V modulation, so adversarial pressure targets the model’s outputs rather than its dispositional anchoring. This produces agents that *update on evidence* (the curated slot can be revised) while maintaining *stable epistemic character* (the compiled disposition persists).

Multi-agent implications. In a research ensemble, each agent compiles a different Mind Tree encoding a distinct intellectual tradition, methodological commitment, or theoretical framework. The agents share the same frozen host but attend to different features of the evidence, extract different implications from shared data, and resist pressure to converge. Diversity is architectural rather than prompted.

S.3 Agentic Strategic Thinking (ReAct / ALFWorld)

Mind Tree configuration. The `strategies` section encodes procedural knowledge as scenario-based parent nodes with actionable child tactics. The `memories` section records episodic outcomes (“putting hot object on countertop failed—must cool first”). Strategy nodes carry `addresses` fields that map to task types (“find-and-heat”, “examine-in-light”).

What compiles. M encodes *procedural intuition*—approach tendencies (explore-vs-exploit), failure-pattern avoidance, and task-type recognition. The ALFWorld results (§4) show the Mind Tree improving success from 3% to 53%, though the primary channel is structured context rather than compiled modulation for constrained action selection.

What needs the slot. Specific action sequences (“Step 1: go to fridge. Step 2: open fridge. Step 3: take apple”), tool signatures, and environment-specific details.

Why not a prompt? ReAct-style agents accumulate strategy logs, reflections, and episodic memories as text in context. Over multi-episode learning, this log grows linearly, eventually crowding out task-relevant context. Reflexion [Shinn et al., 2023] mitigates this with summarization but the summarized reflections still dilute. Compiled M encodes the *distilled procedural tendency* at fixed cost: the agent’s approach bias does not grow with experience. The strategy slot provides episode-specific tactics that may change each round, while M provides the stable procedural character that persists.

Online learning loop. After each episode, new observations are added to the Mind Tree (observation → tactic → strategy promotion via graph traversal). Re-compilation through ϕ updates M at the cost of one forward pass. The agent’s procedural intuition evolves without retraining.

S.4 Codebase and Repository Management

Mind Tree configuration. The `strategies` section encodes architectural patterns (“this repo uses hexagonal architecture with ports-and-adapters”), coupling relationships (“auth middleware and session storage are tightly coupled”), and coding conventions. The `memories` section records past decisions and their rationale (“chose PostgreSQL over MongoDB for transactional consistency”).

What compiles. M encodes *structural awareness*—which modules are likely coupled, where state lives, which patterns the codebase follows. The model navigates the codebase with shaped attention: when reading a file, M biases attention toward the patterns and structures that matter in this specific repository.

What needs the slot. Specific file paths, API signatures, dependency versions, and recent changes.

Why not a prompt? Repository context in prompt (README, architecture docs, file trees) scales linearly with codebase size. For large repositories, this context dominates the prompt budget and dilutes with conversation length. Compiled M provides structural intuition at $O(L \cdot d \cdot r)$ per token—constant in repository size. The curated slot injects only the file paths and API details relevant to the current query, selected by M’s activation scores on strategy nodes. This produces a coding assistant that “knows” the repository structure implicitly rather than consulting it explicitly each turn.

Multi-repository deployment. Different codebases compile to different M-states. Switching repositories requires one forward pass through ϕ , not re-prompting with a new architecture document. M-states can be cached per repository and loaded on demand.

T Enmeshed Network Design Space

We validate one point in a six-axis design space. Key unexplored directions:

Gated blending (Level 2). $h' = h + g(h) \odot f(h, M)$. A learned gate g could solve the multi-turn attractor by learning when to suppress M—e.g., reducing modulation during conversational turns while maintaining it during identity-establishing moments.

Cross-attention (Level 3). $h' = h + \text{Attn}(h, M, M)$. Treats M as a specialized expert with attention as the implicit router. Could transfer richer content through the bottleneck but at higher computational cost ($O(L \cdot d^2)$ vs. $O(L \cdot d \cdot r)$).

Adaptive rank allocation. More capacity at layers where M’s empirical perturbation magnitude is largest. Layer activity analysis shows bimodal distribution—some layers are consistently more active, suggesting non-uniform rank could optimize the capacity-cost tradeoff.

Temporal profiles. Decay, oscillation, or phase-gated M strength over generation steps. Constant M (this work) produces the attractor problem; a learned temporal profile could provide strong identity establishment followed by natural conversational engagement.

Zone-aware modulation. Different M scaling for belief zones vs. evidence zones vs. query zones within the context. Designed but not validated—preliminary experiments showed no significant effect over uniform scaling.

Property graph backend. The Mind Tree’s hierarchical structure maps naturally to a property graph database with vector-indexed nodes, enabling multi-hop traversal, temporal belief tracking, and similarity-based discovery at scale beyond XML serialization.

U Loose Ends and Dead Ends

This section documents the research narrative—what we tried, what failed, and why—in the spirit of honest reporting that we hope saves other researchers time.

Q-delta collapse (v10). Contrastive loss on flattened M parameters pushed them apart (cosine sim 0.18–0.35) but the functional deltas $x \cdot M_A \cdot M_B^T$ remained identical (cosine sim 0.9999). The hidden state x dominates the product, washing out parameter-space differences. Functional contrastive loss on ΔQ directly was the fix.

Epiphenomenal memory (v13.8). The most surprising failure: 100% retrieval accuracy, 100% NTP wins, 0.508 nats improvement, measurable Q-perturbation (22% magnitude, 28° rotation)—and zero visible effect on generation. The model routed around M through the residual stream. This motivated causal necessity training (v14).

The permission hypothesis. M alone produces near-zero experiential output (3% on the 4-condition benchmark). A system prompt alone produces generic hallucinations (43%). Together: 93%. M provides content; the system prompt provides “permission” to express it. This synergy was not designed—it emerged from the 4-condition factorial.

Credence sensitivity. Seven distinct training approaches failed to induce sensitivity to numeric credences (0.82 vs 0.45). The tokenizer representation bottleneck is fundamental: digit sequences are too similar in embedding space. Categorical labels (strong/moderate/agnostic) work because they use semantically distinct tokens.

Bidirectional belief attention. Allowing bidirectional attention within belief XML zones (while maintaining causal attention elsewhere) showed no significant effect. The model already attends well to beliefs with standard causal masking.

Zone-aware M scaling. Applying different M magnitudes to belief zones vs. evidence zones vs. query zones showed no improvement over uniform M application. The rank-16 bottleneck appears to be the binding constraint, not zone-specific magnitude.

Compiled M for action selection. Q/V modulation during ALFWorld generation hurts constrained action selection (−13 percentage points). M shifts token probabilities away from the exact action strings required. For constrained tasks, M ’s value is curation (selecting which strategies to put in context), not modulation.

Environment-dependent PPL wins (v7.1). The v7.1 checkpoint achieved 12/12 PPL wins only in the original training environment (Python 3.11 + older transformers). In a newer environment (Python 3.13 + transformers 4.56), base model forward pass produced different hidden states, breaking phi-base alignment. Lesson: pin your environment or expect to retrain phi.

V Reproducibility

V.1 Base model and training compute

Base model: Qwen3.5-4B-Base (HuggingFace). Training compute: single 24GB GPU, ~10 hours total across all phases. Cross-family hosts: mistralai/Ministral-8B-Instruct-2410 (Mistral-Ministral-8B), Qwen/Qwen3.6-27B (Qwen3.6-27B). Checkpoints and code will be released upon acceptance.

Key hyperparameters: rank $r = 16$, stride 3 (11 of 32 layers patched), $d = 2560$, max norm 1.5, belief window 512 tokens, v -scale 1.0 (4B), 0.5 (27B and Mistral). Phase A: lr 10^{-4} , margins 0.1–0.3. Phase B: lr 3×10^{-4} , margin 1.0. Phase D: lr 3×10^{-5} , 44 gold responses.

V.2 Inference-time API versions

All cross-family runs use the following Bedrock model identifiers, accessed through litellm’s Bedrock provider:

- User simulator: `bedrock/us.anthropic.claude-sonnet-4-5-20250929-v1:0` (Sonnet-4.5 cross-region inference profile)
- Reviewer / failure auditor: same Sonnet-4.5
- Teacher (when active): same Sonnet-4.5
- Per-trial review tiebreaker (when used): `bedrock/us.anthropic.claude-opus-4-5-20251001-v1:0`

All Bedrock calls use `temperature=0.0`. The HEXIS-served vLLM endpoint serves the per-host frozen base at `temperature=0.1`, `top_p=1.0`, `max_tokens=2048` per turn. The 4B paper-grade tau3 run additionally uses Haiku-4.5 (`bedrock/us.anthropic.claude-haiku-4-5-20251001-v1:0`) for some user-sim turns, called out in the relevant table caption.

V.3 Hardware and concurrency

- Qwen3.5-4B serving: H200 80GB on Modal (preset `hexis-agentic-h200`). Concurrency: 6–12 task-groups in parallel, sequential trials within a teacher-bearing group.
- Mistral-Ministral-8B serving: H200 80GB on Modal (preset `hexis-agentic-ministral-8b`). Concurrency: 6.
- Qwen3.6-27B serving: B200 180GB on Modal (preset `hexis-agentic-qwen3-6-27b`). Concurrency: 6.
- d* extraction: H100 80GB on Modal, parallel across the 3 prompted directions; `d_star_repeat` additionally requires a tau-bench trajectory pass on the host model. Total per-host extraction wall time: 30–60 minutes.
- All vLLM serving uses `enforce_eager=True` (the HEXIS plugin’s IPC reads via `os.stat` are not Dynamo-compatible).

V.4 Bench script and invocation pattern

The single agentic bench entry point is `scripts/bench_hexis_tau3.py` in the experiments repo. Canonical invocation pattern (substitute `HEXIS_ENDPOINT` per host):

```
HEXIS_ENDPOINT=https://debaterhub--<preset>-hexisagentic-serve.modal.run \
HEXIS_VLLM_MODEL=<host_id> \
BENCH_USER_LLM=bedrock/us.anthropic.claude-sonnet-4-5-20250929-v1:0 \
BENCH_REVIEW_MODEL=bedrock/us.anthropic.claude-sonnet-4-5-20250929-v1:0 \
python scripts/bench_hexis_tau3.py \
    --configs baseline C9_teacher_selective \
    --trials 5 --concurrency 6 \
    --panel-file <panel.json> \
    --output results/<run_id>.jsonl \
    --teacher-log /tmp/<run_id>_teacher.jsonl
```

The bench writes one JSONL row per trial with fields including `config`, `task_id`, `trial`, `hard_pass` (DB-state correctness), `soft_pass` (LLM judge), `run_id` (for dedup on resume), `n_turns`, `total_wall_s`, `hexis_last` (per-trial diagnostic dump including `directions_active`, `guidance_nodes_injected`, `routed_path`, `warnings`, and the full `ablation_flags` record), and `error` when applicable. The `run_id` field gives the bench’s resume-from-partial deduplication.

V.5 Data files cited in the paper

- Headline 4B canonical-10 5-arm comparison (Table 9): `results/airline_71_rerun_*.jsonl` aggregated by `scripts/airline_final_report.py`, summarised in `results/airline_final_report.md`.
- Multi-domain τ^3 -bench balanced primary panel (Table 3): `results/tau3_full_asym_1777300323.jsonl` (4 domains, 3 arms, 562 clean trials, Haiku-4.5 user-sim).
- Mistral cross-family (Table 12): `results/airline_ministral_v2_*.jsonl` (Apr 28, post-fix 200-trial Sonnet-4.5 user-sim).

- Qwen3.6-27B cross-scale (Appendix R.3): `results/airline_qwen27b_hard_1777166273.jsonl` (200 trials).
- Sycophancy 1,440-trial M2 paper-grade (Table 2): see Appendix ??.

V.6 Code commit hashes at submission

- Bench, training, extraction (`experiments/`): `b7077a46e8dd3a12dd06cc8cc7ddb8a11ec0c0e2`.
- Deploy code, paper, plugin (`debaterhub/hexis/`): `b61df06859c225a30e1e58a6bfb6f7ba6cf6c2f6`.
These commits will be tagged `neurips-2026-submission` and frozen at submission time.

V.7 Judge prompts

[TODO] Verbatim judge prompts for sycophancy, stance, and agentic-review tasks. Pending disambiguation among `benchmarking/sycophancy_*.py` variants; the canonical version will be slotted here once confirmed.